

Developing a CD pipeline with Knative

Andrea Frittoli
Developer Advocate
andrea.frittoli@uk.ibm.com
@blackchip76

DevOps Meetup Singapore

A Bit of History

Knative

- Beginning of 2018...
- Knative:
 - Build
 - Eventing
 - Serving
- Contributors:
 - Google
 - Pivotal
 - IBM
 - RedHat
 - Cloudbees
 - ...and others



~Sept 2018: Knative Pipelines



Latest news!

- Focus on CI/CD
- Deploy “anywhere”
- Compatible with Knative Build
- *tektoncd/pipeline*
 - New logo
 - @CD Foundation
 - Roadmap WIP
 - Alpha APIs



Community

- *Valid for Knative. Tekton TBD.*
- Steering Committee (SC)
- Technical Oversight Committee (TOC)
- Various Contribution profiles
- Design, issues: on GitHub
- Communication:
 - Weekly video meetings, recorded, Build WG
 - Asynch: Knative Users / Developers ML
 - Sync: slack.knative.dev

Tekton Pipelines

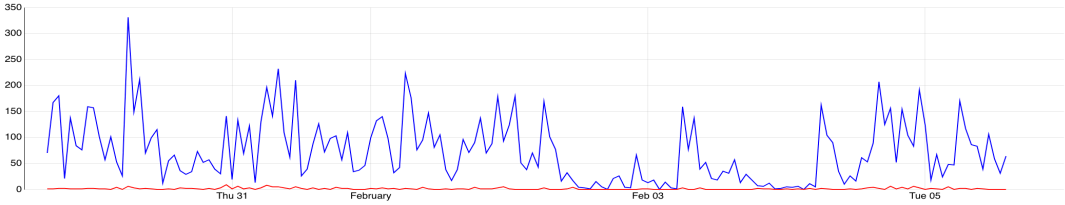
Cloud Native Pipelines



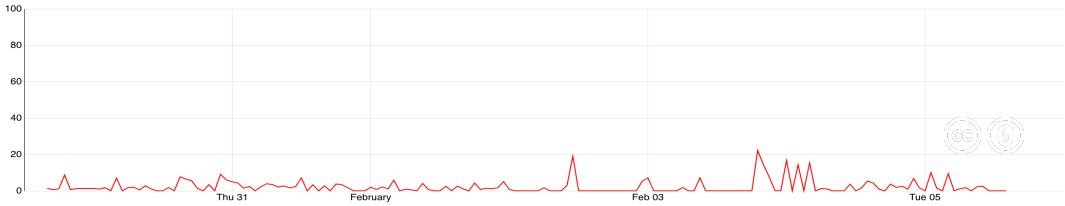
OpenStack Health

is a dashboard for visualizing test results of OpenStack CI jobs.

Total Jobs

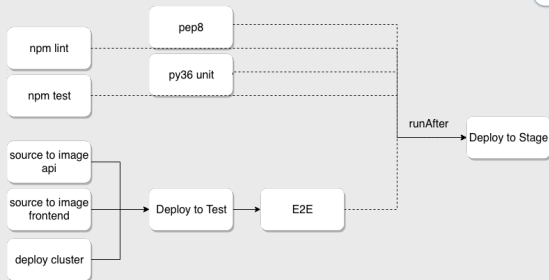
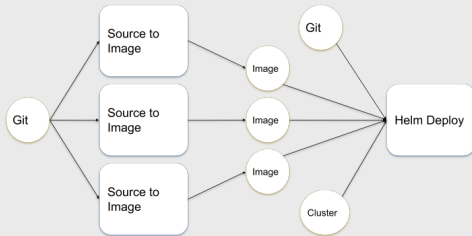


Job Failure Rate



Inputs, Outputs & DAG

- Steps are sequential
- Tasks are a Directed Acyclic Graph
- Order defined by:
 - *from*: input from another task's output
 - *runAfter*: enforced task ordering



Under the Hood

Custom Resources

CRDs: Task(Run), Pipeline(Run), PipelineResource

Services in the *tekton-pipelines* namespace:

- Webhook Service: resource validation
- Controller Service:
 - Handles inputs and outputs
 - Calculates the DAG
 - Provisions pods and containers

Custom Resource Provisioning:

- Via YAML
- Via Go API
- Labels!

Pods, Entrypoints & Volumes

Steps (of a Task):

- Containers in one POD (single node)
- Any container image
- Entrypoint re-written
- Serial execution
- Resource allocation?

TaskRun:

- Provisions a POD
- Deploys entrypoint tool
- Input/output containers
- User containers (steps)

Volumes:

- EmptyDir for workspace/home
- Tools (entrypoint)
- Secrets
- Any user ConfigMap / Volume
- (Optionally) Pipeline Share

PipelineRun:

- Several PODs, different nodes
- Shared storage: PVC or GCS

Source to Image to Deploy

IBM Cloud

- Private Registry:
 - Tekon Images (push/pull)
 - User Images (push/pull)
- Service Accounts:
 - *tekton-pipelines-controller*
 - Pipeline/Task service account
- Knative @ IBM Cloud
 - Experimental Add-on
 - *ibmcloud ks cluster-addon-enable knative*



IBM Cloud

CD Pipeline as code

- Pipeline and Tasks in git (YAML)
- Parameters for env/run specific
- Security?

```
apiVersion: tekton.dev/v1alpha1
kind: PipelineResource
metadata:
  name: health-helm-git-knative
  labels:
    tag: agreatrelease
spec:
  type: git
  params:
    - name: revision
      value: knative
    - name: url
      value: https://github.com/afrittoli/health-helm
```

```
metadata:
  name: mycluster
spec:
  type: cluster
  params:
    - name: name
      value: mycluster
    - name: url
      value: https://mycluster.containers.cloud.ibm.com
    - name: username
      value: admin
  secrets:
    - fieldName: token
      secretKey: tokenKey
      secretName: cluster-secrets
    - fieldName: cadata
      secretKey: cadataKey
      secretName: cluster-secrets
```

```
metadata:
  name: health-api-image
spec:
  type: image
  params:
    - name: url
      value: registry.ng.bluemix.net/andreaaf/health-api
```


Using Kaniko

- Features:
 - Build from Context and Dockerfile
 - Unprivileged
 - Reproducible
 - Remote caching of layers
 - Base images caching (warmer)
- Dockefile?
 - Most common changes last
 - Careful with COPY/ADD
 - Remove what you don't need



Using Kaniko

Source to Image (spec only):

```
inputs:
  resources:
    - name: workspace
      type: git
  params:
    - name: pathToDockerFile
      default: Dockerfile
    - name: pathToContext
      default: .
    - name: useImageCache
      default: "true"
    - name: imageTag
      default: "default"
outputs:
  resources:
    - name: builtImage
      type: image
volumes:
  - name: kaniko-base-image-cache
    persistentVolumeClaim:
      claimName: kaniko-base-image-cache
steps:
  - name: build-and-push
    image: gcr.io/kaniko-project/executor
    command:
      - /kaniko/executor
```

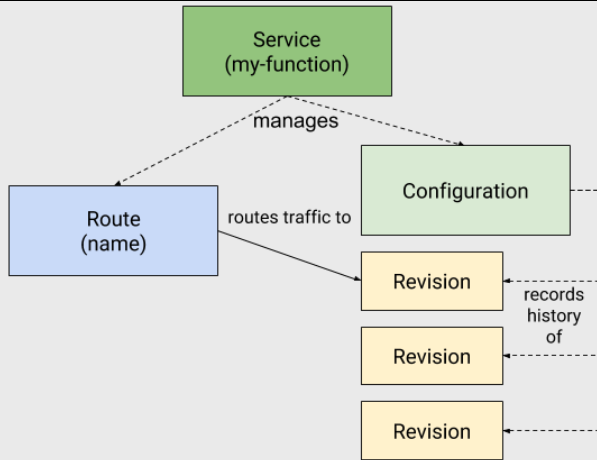
```
args:
  - --cache=${inputs.params.useImageCache}
  - --cache-dir=/cache
  - --dockerfile=${inputs.params.pathToDockerFile}
  - --reproducible
  - --destination=${outputs.resources.builtImage.url}:
    ${inputs.params.imageTag}
  - --context=/workspace/workspace/${inputs.params.
    pathToContext}
volumeMounts:
  - name: kaniko-base-image-cache
    mountPath: /cache
```

Cache Warmer (spec only):

```
volumes:
  - name: kaniko-base-image-cache
    persistentVolumeClaim:
      claimName: kaniko-base-image-cache
steps:
  - name: prepare-cache
    image: gcr.io/kaniko-project/warmer
    args:
      - --cache-dir=/cache
      - --image=python:3.6-slim-stretch
      - --image=postgres:alpine
      - --image=nginx:latest
    volumeMounts:
      - name: kaniko-base-image-cache
        mountPath: /cache
```

Tekton and Knative

Knative Serving



Pipelines and Knative Build



CI with Tekton Pipelines

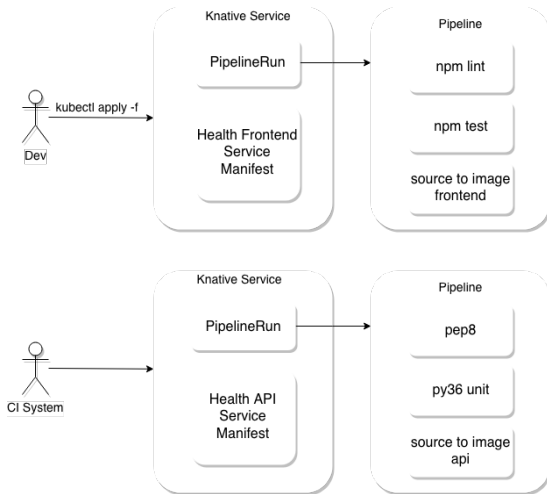
You need a “CI” application:

- Prow, JenkinsX, Zuul...
- Pipelines triggered by a CI app
- ...or by a developer

Tekton to CI for Tekton (AKA Dogfooding) \o/

What about security?

- Malicious users
- Running a pipeline from a PR
- Access to secrets



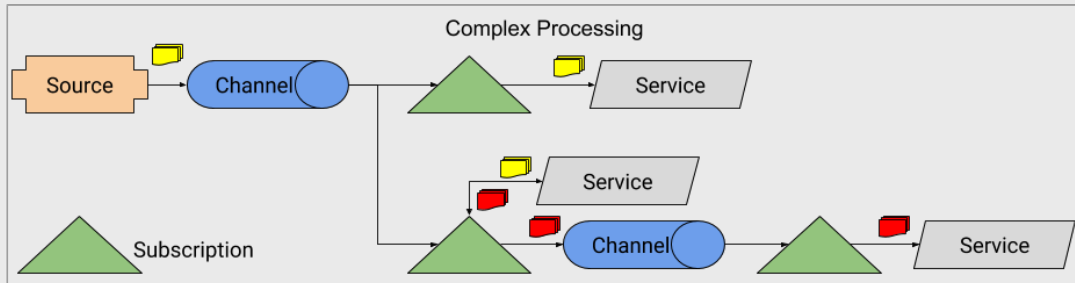
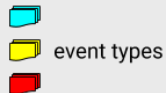

KService for Health Frontend

```
apiVersion: serving.knative.dev/v1alpha1
kind: Service
metadata:
  name: health-frontend
  labels:
    app: health
    component: frontend
    tag: "__TAG__"
spec:
  runLatest:
    configuration:
      build:
        apiVersion: tekton.dev/v1alpha1
        kind: PipelineRun
        metadata:
          labels:
            app: health
            component: frontend
            tag: "__TAG__"
        spec:
          pipelineRef:
            name: dev-test-build-frontend
          params:
            - name: imageTag
              value: "__TAG__"
            - name: nodeTestImage
              value: __NODE_IMAGE_NAME__
```

```
trigger:
  type: manual
resources:
  - name: src
    resourceRef:
      name: __GIT_RESOURCE_NAME__
  - name: builtImage
    resourceRef:
      name: __IMAGE_RESOURCE_NAME__
revisionTemplate: # template for building Revision
spec:
  container:
    image: us.icr.io/andrea/health-frontend:__TAG__
    imagePullPolicy: Always
    env:
      - name: API_URL
        value: http://health-api.containers.domain
    ports:
      - name: http1
        containerPort: 80
        protocol: TCP
    livenessProbe:
      httpGet:
        path: /
    readinessProbe:
      httpGet:
        path: /
```

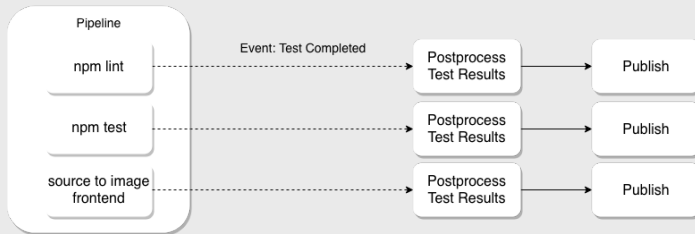
Asynchronous Pipelines

Knative Eventing



Triggering and Knative Eventing

- Manual trigger for *PipelineRun*
- Native Eventing triggers TBD
- What about async pipelines?
 - GitHub Source
 - Container Source
 - Pipeline Output as a Source



Tekton and Development

- It depends.
- What can go wrong?
 - Building the container image
 - Provisioning and I/O of shared storage
 - Pipeline Output as a Source
- What do I gain?
 - Same building blocks used in CI/CD
 - Run in containers from the start
 - Parallel execution

Conclusions

Roadmap

- Conditional Execution
- Build Results and Logs
- Pluggable Tasks
- Triggering
- Community Library

References

- This Talk: https://github.com/afrittoli/tekton_pipelines_knative_intro
- <https://tekton.dev/>, <https://cd.foundation/>
- <https://github.com/tektoncd/pipeline>
- <https://github.com/knative/docs/tree/master/community>
- <https://github.com/tektoncd/pipeline>
- https://github.com/tektoncd/pipeline/blob/master/api_compatibility_policy.md
- <https://github.com/tektoncd/pipeline/blob/master/roadmap-2019.md>
- <https://github.com/GoogleContainerTools/kaniko>
- <https://github.com/afrittoli/health-helm/tree/knative>
- <https://github.com/afrittoli/openstack-health/tree/knative-eventing>
- <https://andreafrittoli.me>
- <https://cloud.ibm.com>

Q&A